CROSS-AD · OPTION 1 PROTOTYPE

Shared EEG–MRI embedding from unpaired cohorts: results

22 October 2026 · code `train_option1.py` (job script `train_option1.qsub`) · outputs `results/option1/`

Bottom line. Training EEG together with unpaired ADNI MRI raised EEG AUC over an EEG-only network (0.691 → 0.725), but the ablations show the gain is not MRI information: shuffling the MRI labels keeps most of it, and replacing MRI with age/sex alone does slightly better than MRI (-0.012). The joint model only matches plain EEG logistic regression (-0.002). The contrastive loss acts as a regularizer for a small network; we find no evidence that unpaired MRI improves EEG prediction (or vice versa).

1. Data

Modality	Cohort	Subjects (AD / CN)	Input features
EEG	BrainLat (2 sites) + OpenNeuro ds004504 (1 site), resting state	71 / 61 (35/32 + 36/29)	26 spectral features: relative δ, θ, α, β power (global + 4 scalp regions), log θ/α ratio (global + regions), peak alpha frequency
MRI	ADNI, baseline 3T, CN age/sex-matched 2:1	294 / 576	321 UCSF FreeSurfer 7 regional measures (volume, area, thickness); volumes ÷ intracranial volume; 2 near-empty columns dropped

No subject has both modalities. EEG was preprocessed identically for both datasets (average reference, 1–30 Hz, 256 Hz, ≤300 s, Welch PSD). Age, sex and education are *not* model inputs because they are confounded with diagnosis in BrainLat. Details: `reports/data_introduction.pdf`.

2. Method

Model

- **Encoders**: one MLP per modality (EEG 26→32→16; MRI 321→64→16; ReLU, dropout 0.3) into a shared 16-d embedding.
- **Shared head**: one linear AD/CN classifier used by both modalities, so "AD" means the same direction in the embedding.
- **Cross-modal supervised contrastive loss** ($\tau = 0.1$): each EEG embedding is pulled towards MRI embeddings with the same label and away from the other label, and vice versa. No subject pairing is used — only labels.
- **Loss** = $BCE_{target} + 1.0 \cdot BCE_{auxiliary} + 0.5 \cdot$contrastive; BCE class-weighted for imbalance.
- Full-batch Adam, lr 0.001, weight decay 0.001, 300 epochs. All hyperparameters fixed before running; nothing was tuned on test folds.
- At test time a subject is predicted **from its own modality only**.

Evaluation

- **EEG CV**: 5-fold × 10 repeats over all EEG subjects, stratified by dataset × diagnosis; all MRI subjects are auxiliary training data.
- **EEG transfer**: train on one EEG dataset (+ MRI), test on the other; 10 seeds.
- **MRI CV**: 5-fold × 3 repeats over ADNI; all EEG subjects auxiliary.
- Imputation and z-scoring fitted on training rows only (EEG per dataset; the unseen transfer dataset is z-scored with its own, label-free statistics).
- Metric: AUC of out-of-fold predictions per repeat; every model uses the same splits, so differences are paired.
- **Comparisons**: logistic regression and the same network without the other modality; ablations without the contrastive loss, with auxiliary labels shuffled, and with auxiliary nodes that carry only age + sex.

3. Results

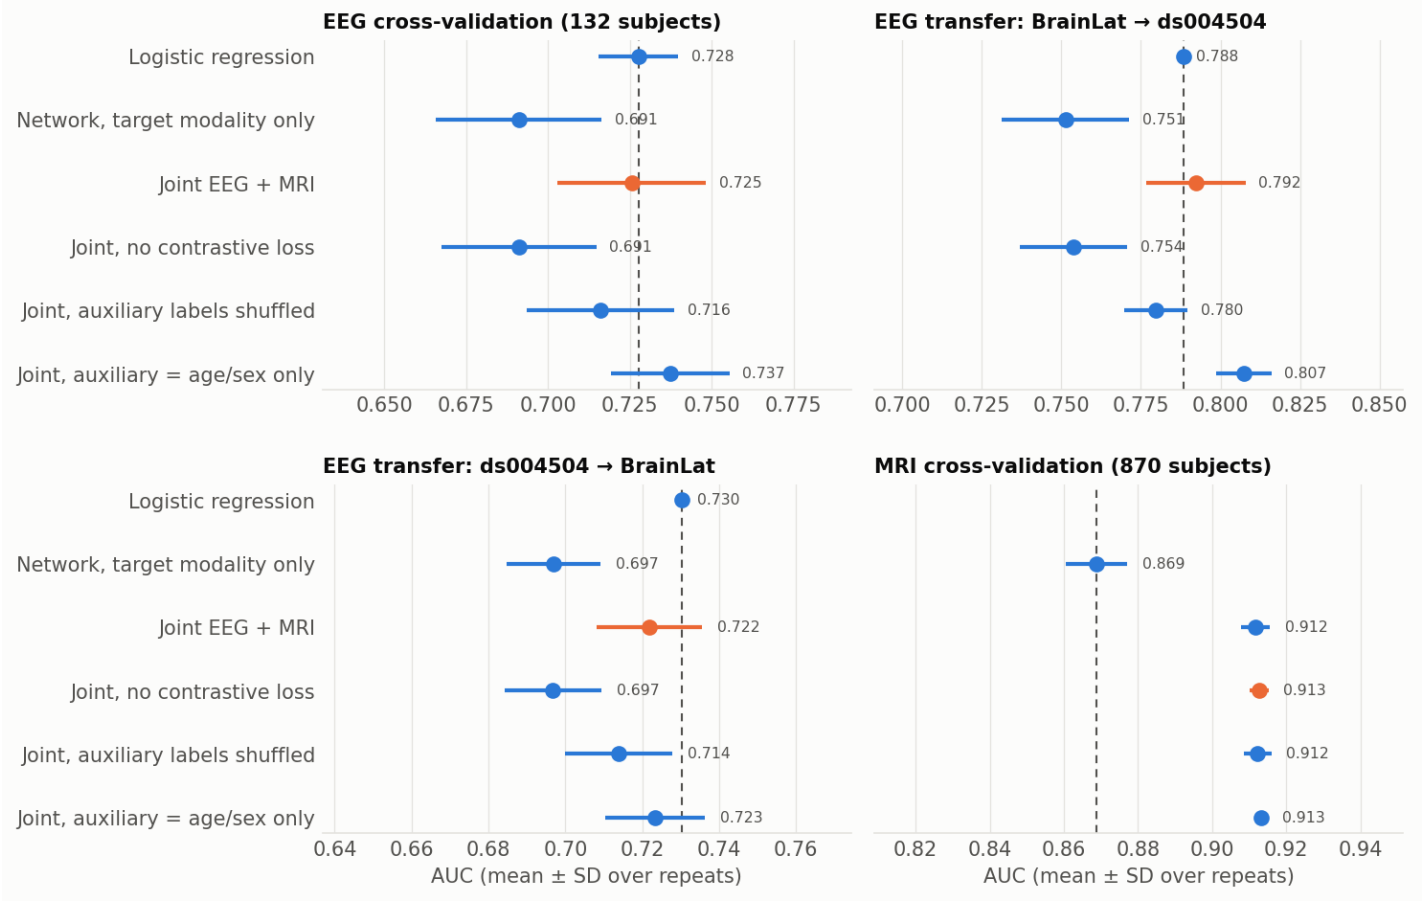


Figure 1. AUC per model (mean ± SD over repeats). Orange = joint EEG + MRI model; dashed line = logistic regression.

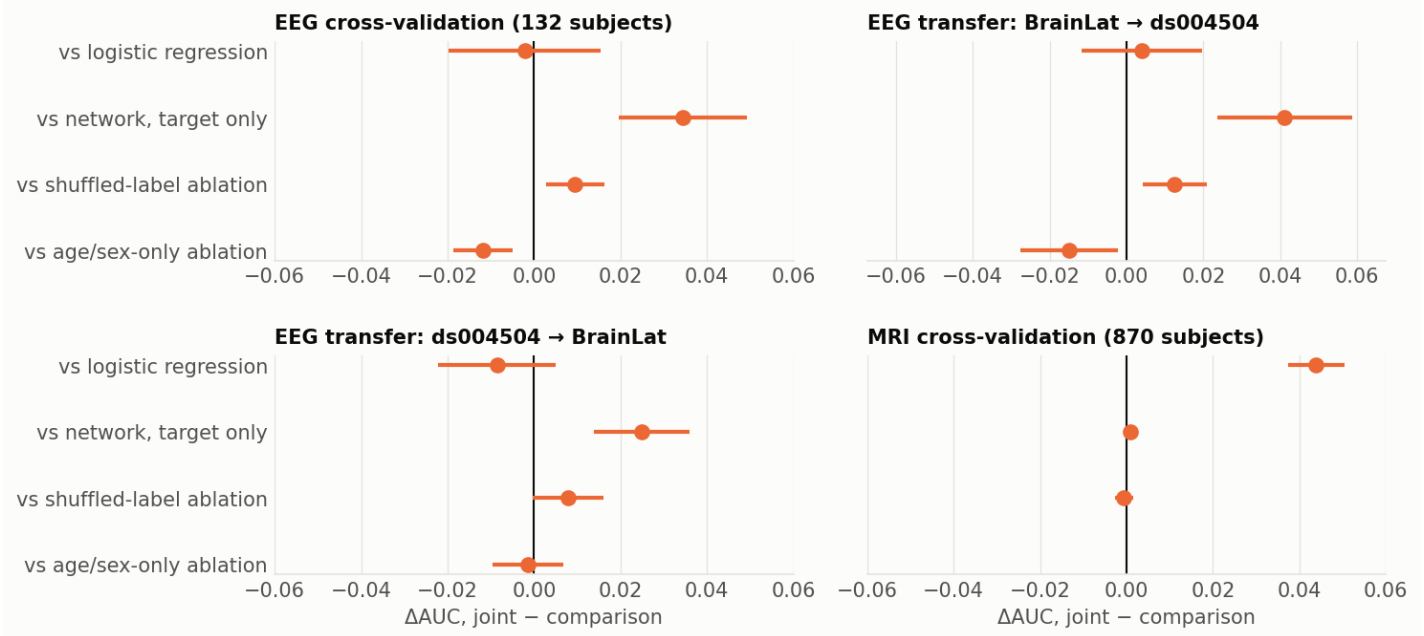


Figure 2. Paired AUC difference of the joint model against each comparison (same splits). Right of 0 = joint better. If MRI carried useful information, the joint model should beat the shuffled and age/sex-only ablations clearly; it does not.

Joint model vs comparisons (paired)

Track	Joint model	ΔAUC (mean ± SD)	Joint better in
EEG cross-validation (132 subjects)	vs logistic regression	-0.002 ± 0.018	3 / 10
EEG cross-validation (132 subjects)	vs network, target only	+0.034 ± 0.015	10 / 10
EEG cross-validation (132 subjects)	vs shuffled-label ablation	+0.009 ± 0.007	9 / 10
EEG cross-validation (132 subjects)	vs age/sex-only ablation	-0.012 ± 0.007	0 / 10
EEG transfer: BrainLat → ds004504	vs logistic regression	+0.004 ± 0.016	4 / 10
EEG transfer: BrainLat → ds004504	vs network, target only	+0.041 ± 0.018	10 / 10
EEG transfer: BrainLat → ds004504	vs shuffled-label ablation	+0.013 ± 0.008	10 / 10
EEG transfer: BrainLat → ds004504	vs age/sex-only ablation	-0.015 ± 0.013	1 / 10
EEG transfer: ds004504 → BrainLat	vs logistic regression	-0.009 ± 0.014	2 / 10
EEG transfer: ds004504 → BrainLat	vs network, target only	+0.025 ± 0.011	10 / 10
EEG transfer: ds004504 → BrainLat	vs shuffled-label ablation	+0.008 ± 0.008	7 / 10
EEG transfer: ds004504 → BrainLat	vs age/sex-only ablation	-0.002 ± 0.008	4 / 10
MRI cross-validation (870 subjects)	vs logistic regression	+0.044 ± 0.007	3 / 3
MRI cross-validation (870 subjects)	vs network, target only	+0.001 ± 0.001	2 / 3
MRI cross-validation (870 subjects)	vs shuffled-label ablation	-0.001 ± 0.002	1 / 3

Per-model AUC for every track is in Figure 1 and `results/option1/summary.csv`.

EEG CV AUC within each EEG dataset

subset	BrainLat	ds004504
Model		
Logistic regression	0.653	0.798
Network, target modality only	0.613	0.774
Joint EEG + MRI	0.645	0.804
Joint, no contrastive loss	0.612	0.774
Joint, auxiliary labels shuffled	0.634	0.796
Joint, auxiliary = age/sex only	0.661	0.808

4. Interpretation

- **EEG** separates AD from CN with AUC ≈ 0.73 overall (≈0.80 within ds004504, ≈0.65 within BrainLat) and transfers across datasets at 0.73–0.79.
- **MRI** regional features reach AUC ≈ 0.91; adding EEG as auxiliary data changes nothing (+0.001).
- The joint model's gain over the EEG-only network (+0.034) disappears against the ablations: the shared head alone adds 0.000, the shuffled-label ablation keeps most of the gain, and age/sex-only auxiliary nodes do better than MRI. The contrastive term regularizes the small EEG network, which otherwise overfits; logistic regression gets the same benefit for free.
- This is the controlled negative result anticipated in the proposal: with unpaired cohorts, the label-aligned shared embedding does not transfer disease information between modalities beyond regularization.

5. Limitations and next steps

- Hyperparameters were fixed, not tuned (nested CV would tune them without test leakage); the EEG-only network had no early stopping, so logistic regression is the fairer baseline.
- 132 EEG subjects: AUC differences of ≈0.01 are within noise (SD across repeats 0.01–0.03).
- BrainLat confounds (recording type matches diagnosis, AD older and less educated) limit what the BrainLat numbers mean.
- Next: nested-CV tuning, EEG connectivity features, global-only EEG features for cross-dataset robustness, feature importance; and paired EEG–MRI fusion once BrainLat MRI download access is granted.

Acknowledgments and data sources

Data collection and sharing for this project was funded by the Alzheimer's Disease Neuroimaging Initiative (ADNI) (National Institutes of Health Grant U01 AG024904) and DOD ADNI (Department of Defense award number W81XWH-12-2-0012). Data used in preparation of this work were obtained from the ADNI database (adni.loni.usc.edu). As such, the investigators within the ADNI contributed to the design and implementation of ADNI and/or provided data but did not participate in analysis or writing of this report. A complete listing of ADNI investigators can be found at http://adni.loni.usc.edu/wp-content/uploads/how_to_apply/ADNI_Acknowledgement_List.pdf

BrainLat EEG: Prado P. et al., The BrainLat project, a multimodal neuroimaging dataset of neurodegeneration from underrepresented backgrounds, Scientific Data 10, 889 (2023); Synapse syn51549340.

ds004504 EEG: Miltiadous A. et al., A dataset of scalp EEG recordings of Alzheimer's disease, frontotemporal dementia and healthy subjects from routine EEG, Data 8(6), 95 (2023); OpenNeuro ds004504 (CC0).